

Playing checkers using reinforcement learning

An investigation into playing checkers using a reinforcement learning via a neural network library, with an implementation of automatic differentiation, built from scratch.

Edgar Maddocks

Contents

1 Analysis	1
1.1 Background	1
1.2 Evidence of Analysis	2
1.3 Modelling of the problem	10
1.4 Set of objectives	19
2 Design	21
2.1 File Structure and Project Organisation	21
2.2 Checkers	23
2.3 Autodifferentiation Engine	27
2.4 Neural Network Library	28
2.5 Reinforcement Learning Solution	29
3 Research Log	31
3.1 Checkers	31
3.2 Neural Networks	31
3.3 Automatic Differentiation	31
3.4 AlphaZero	31
3.5 Existing Auto differentiation systems	31
3.6 Reinforcement Learning	31

1 Analysis

1.1 Background

This project is an investigation into the use of reinforcement learning to play games (in this case checkers/checkers). The model will use self-play and Monte Carlo tree search algorithms, coupled with a multi-headed neural network to understand the game and estimate optimal actions.

The neural network will be built using a library, that has an autograd engine implemented, all built from scratch. This investigation will however use some scientific computing libraries for faster simple matrix operations (such as numpy). Some more complex functions (such as cross-correlation and its derivative) will also be built from scratch.

1.2 Evidence of Analysis

1.2.1 Current systems

1.2.1.1 Checkers

checkers is an English board game played on an 8x8 checkered board, identical to a chessboard. Each player begins a game with 12 pieces, usually flat round discs. The pieces and board are usually black and white, and will be referred to as such. The board is first placed between the two players such that the bottom right-hand corner is a white square, for both players.

A coin is tossed to decide who plays black, and that player has the first move. Each player places their pieces on the 12 black squares closest to themselves. The setup of the board can be seen in Figure 1.

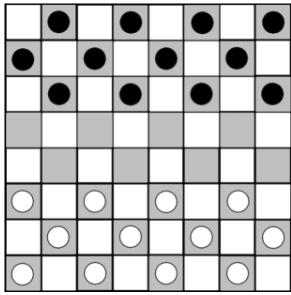


Figure 1: An image showing the starting position of a game of checkers

can be seen in Figure 3.

Unlike a regular move, a capturing move may make more than one 'hop'. This is if the capture places the piece in a position where another capture is possible. In this case, the additional capture must be made. The capture sequence can only be made by one piece per move. i.e. You cannot make one capture with one piece, and then another capture with another piece in the same move.

However, if more than one piece can capture, the player has free choice over which piece to move. Likewise, if one piece can capture in multiple directions then the player has the choice in which direction to move.

Note: it is not compulsory for the player to move in the direction, or with the piece, that will lead to the greatest number of captures in that move.

The pieces only move diagonally (so will always be on black squares) and the aim is to take all of the opposing players pieces, or to put the opposing player in a position with no possible moves. Players take turns moving their shade of pieces. If at any point of the game, a player's piece reaches the opposing players edge of the board, the piece becomes a 'King', and another piece should be placed on top of said piece to indicate so. Unless a piece is crowned and a 'King' it may only move and take pieces diagonally forwards. Kings may move and take both forwards and backwards.

If an adjacent square has an opponents piece and the square immediately beyond the oppositions piece is empty, the opponents piece may be captured. If the player whose go it is, has the opportunity to capture one or more pieces, then they must do so. A piece is taken by moving your own piece over the opposing player's, into the vacant square, and then removing the opposing piece from the board. An example of this process

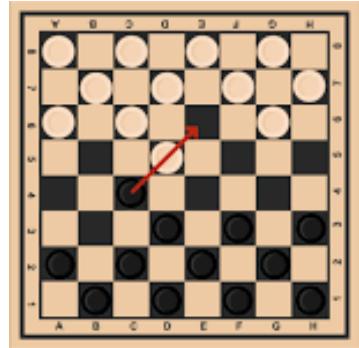


Figure 2: Example of a piece being taken in checkers

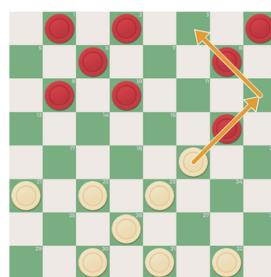


Figure 3: Visualization of multiple captures in one move

A move may only end when the position has no more captures available or an uncrowned piece becomes a King. The game ends when all of a players piece's have been captured, or a player has no available moves.

1.2.1.2 Neural Networks

What is a neural Network

A neural network is a machine learning model which aims to mimic the processes of the human brain. Each network contains inputs and outputs, as well as one or more layers of hidden nodes - which act as artificial neurons.

Neural networks are a supervised learning model, meaning that they learn from labeled data (which has the objective correct answer in the data). They are sometimes referred to as artificial neural networks (ANNs) or simulated neural networks (SNNs).

In a fully connected network, each node is connected once to each node in the next layer - an example of how one node connects to the next layer can be seen below.

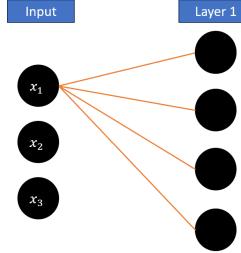


Figure 4: Example of a fully connected node and layer

Neural networks can be modelled as a collection linear regression units.

Linear Regression Unit

A single linear regression unit output has the formula:

$$\hat{y} = \sum_{i=0}^n w_i x_i + b \quad (1)$$

Where \hat{y} is the predicted output, n is the number of inputs, x_i is the i th input, w_i is the weight of x_i , and b is a bias. If, for example, there were 3 inputs the full equation for \hat{y} would be:

$$\hat{y} = w_1 x_1 + w_2 x_2 + w_3 x_3 + b \quad (2)$$

Vectorization of processes

This calculation can be vectorized to improve efficiency and would be notated:

$$\hat{y} = XW + b \quad (3)$$

Where we let

$$W = \begin{bmatrix} w_1 \\ w_2 \\ \vdots \\ w_n \end{bmatrix} \quad (4)$$

$$X = [x_1 \ x_2 \ \dots \ x_n] \quad (5)$$

X here is a row vector as this is the most common format for data as an input to a network (e.g. being read from a csv). This parallelized computation is much faster, and can be parallelized using the GPU to further improve speed and efficiency.

Forward Pass of Dense Layer

In the case of neural networks, lots of these linear regression units can be combined to form a vector of outputs. Each of these regression units will have the same inputs, therefore X can have the same definition. However, W will now be composed of multiple vectors of weights, instead of just one.

Here we let

$$W = \begin{bmatrix} w_{11} & w_{21} & \dots & w_{j1} \\ w_{12} & w_{22} & \dots & w_{j2} \\ \vdots & \vdots & \ddots & \vdots \\ w_{1n} & w_{2n} & \dots & w_{jn} \end{bmatrix} \quad (6)$$

Where j is now the number of nodes in the layer. If we rewrite our forward pass equation to use this weight matrix, with each node as its own regression unit:

$$Y = XW + B \quad (7)$$

Where:

$$Y = [\hat{y}_1 \quad \hat{y}_2 \quad \dots \quad \hat{y}_n] \quad B = [b_1 \quad b_2 \quad \dots \quad b_j] \quad (8)$$

With W and X having the same definition as most recently defined.

Forward Pass of a Convolutional Layer

The main function of a convolutional layer is to extract key features and they are often used in image recognition and classification solutions.

A convolutional layer takes in a 3 dimensional block of data. This can be imagined as the first 2 dimensions being a grid of pixels, and the third dimension being the colour channels for each grid. The trainable parameters in a convolutional network are the kernels/filters and the biases. The kernels can be of any size but must have the same depth as the input data. Each kernel also has a bias matrix, which has the same shape as the output.

After the forward pass, the output will also be 3 dimensional, but here the depth is equal to the number of kernels, and the size is equal to

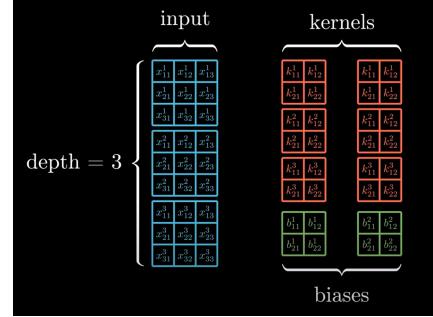


Figure 5: Visualization of a Convolutional layer

$$I - K + 1 \quad (9)$$

Where I is the size of the input grid and K is the size of the kernel grids.

The forward pass can then be simplified, with each output equating to the bias for that kernel plus the sum of the sum of cross correlations between a single channel of the input and its respective channel in the kernel. This can be seen in the figure below, where d is equal to the number of kernels, and n is the number of channels in the input data.

$$\begin{aligned}
Y_1 &= B_1 + X_1 \star K_{11} + \cdots + X_n \star K_{1n} \\
Y_2 &= B_2 + X_1 \star K_{21} + \cdots + X_n \star K_{2n} \\
&\vdots \\
Y_d &= B_d + X_1 \star K_{d1} + \cdots + X_n \star K_{dn}
\end{aligned}$$

Figure 6: Representation of the forward pass of a convolutional layer

This can be written more mathematically using a summation.

$$Y_i = B_i + \sum_{j=1}^n X_j \star K_{ij}, \quad i = 1 \dots d \quad (10)$$

Gradient Descent

Now, to update the weights of our model, we can compute these values in closed form, however, it comes with a large time complexity (greater than $O(n^3)$), and therefore a process called gradient descent is usually employed.

Gradient descent works to minimize the error of a model by iteratively locating a minimum in the error function. For example, a common cost (error) function is mean squared error.

$$MSE(Y, \hat{Y}) = \frac{1}{n} \sum_{i=0}^n (y_i - \hat{y}_i)^2 \quad (11)$$

This function effectively calculates the absolute distance between the predicted value of one data point (\hat{y}_i) and the true value at that same point (y_i), for every value in the dataset (of size n) and takes the mean of these distances. Again, this calculation can be vectorized using the equation below.

$$MSE(Y, \hat{Y}) = \frac{1}{n} ((Y - \hat{Y})^T \cdot (Y - \hat{Y})) \quad (12)$$

If we plot the MSE curve of different weights and biases - of a single regression unit, it looks as follows.

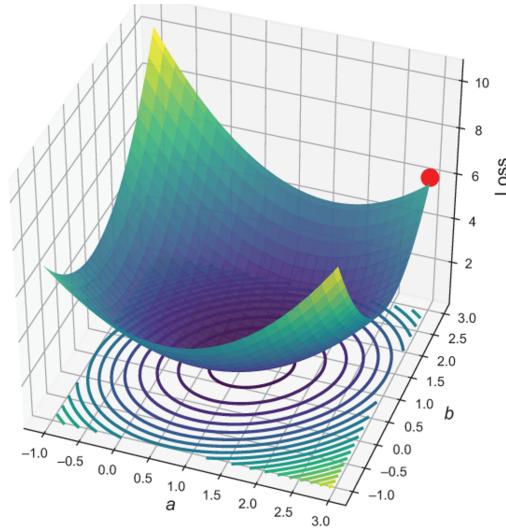


Figure 7: Plot of MSE in a single regression unit

Here the weight has been labelled a , but the bias has remained as b . As we can see, there is a clear local minimum of this error function, and the values that obtain this minimum, is what gradient descent aims to achieve.

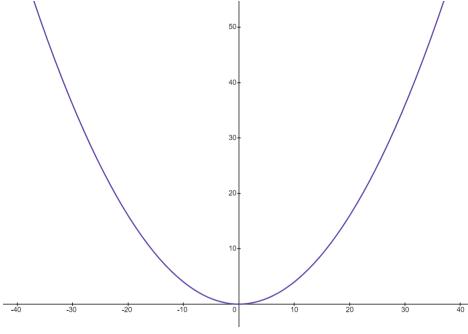


Figure 8: MSE with respect to one parameter

In figure 6 we can see a representation of the MSE plot with respect to only one of the parameters. This makes understanding the process of gradient descent much simpler. For example, the minimum of this function ($f(x) = \frac{x^2}{25}$) is at $x = 0$. So, if we were to have an initial x value of say -40 , and we calculated the gradient to be 3.2 , we could then subtract this from 40 . Resulting in a new value of 36.8 , which is closer to our optimal value of 0 . This process is completed iteratively until the minimum, or near to it, is reached.

To denote this mathematically, we can say that

$$w_{t+1} = w_t - \alpha \frac{\partial E}{\partial w} \quad (13)$$

$$b_{t+1} = b_t - \alpha \frac{\partial E}{\partial b} \quad (14)$$

Where w_{t+1} and b_{t+1} are the parameters at the next timestep, w_t and b_t is the value of the parameters at the current timestep, $\frac{\partial E}{\partial w}$ and $\frac{\partial E}{\partial b}$ are the derivatives of the error with respect to each parameter, and α is the learning rate. This learning rate controls the size of our 'jumps' and prevents the process from beginning to spiral away from the optimal values.

That is a basic overview of gradient descent in a single regression unit, but in a dense network, the process is almost identical. An error function is evaluated, the derivative with respect to the inputs calculated, and the new values updated. The main expense computation-wise is calculating all the derivatives, with respect to each set of parameters, as this requires passing the output from the error function backwards through all the layers and processes. To combat this, automatic differentiation can be used, which tracks computation dynamically at run time to compute derivatives.

1.2.1.3 Automatic Differentiation

As mentioned, automatic differentiation tracks computation dynamically and then computes derivatives using a computational graph. This allows models to have much more complex forward passes, including decision branches and loops where the length of the loop is decided at runtime.

Automatic differentiation operates by differentiating a complex function (that we don't know the derivative of), by treating it as a composition of more elementary functions (of which we do know the derivatives). Additionally, it treats each of these elementary functions as though their output is an intermediate variable when computing the complex function. This becomes very useful when there are multiple inputs to the function, and we only want the derivative to one of these variables. Finally, this process makes use of the chain rule to compute the derivative with respect to the inputs, as will be shown later on.

Firstly, I will walk through an example of the process of automatic differentiation.
Say we have a function:

$$f(x) = \ln \left(\frac{\sin(x_1)}{\cos(x_2)} \right) \quad (15)$$

This function has two arbitrary inputs (perhaps in the case of a neural network these could be our parameters or outputs for example). At first, this is quite a complex function to differentiate but, we can simplify it using intermediate variables. For example

$$v_1 = \sin(x_1) \quad (16)$$

$$v_2 = \cos(x_2) \quad (17)$$

$$v_3 = \frac{v_1}{v_2} \quad (18)$$

$$v_4 = \ln(v_3) \quad (19)$$

$$y = v_4 \quad (20)$$

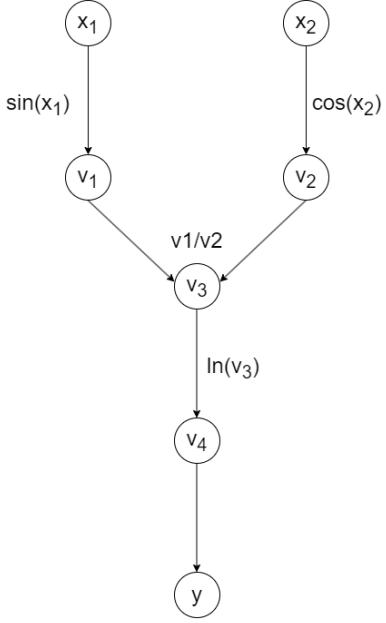


Figure 9: Example Computational Graph

Now that we have declared our intermediate variables, we can see how automatic differentiation splits up the original function into elementary ones. This then allows us to easily take derivatives. Firstly, we can draw out the computational graph for this function, using our intermediate variables, and this can be seen in Figure 7. This graph visualizes the forward pass of the function, and how the intermediate variables compose to form our original function.

Now, when we want to take the derivative with respect to one of these inputs, we can simply reverse through the graph, taking the gradient of each intermediate variable, and using the chain rule to give us our final derivative.

This process would look as follows:

$$\frac{\partial y}{\partial x_1} = \frac{\partial y}{\partial v_4} \cdot \frac{\partial v_4}{\partial v_3} \cdot \frac{\partial v_3}{\partial v_1} \cdot \frac{\partial v_1}{\partial x_1} \quad (21)$$

If we computed each of these derivatives:

$$\frac{\partial y}{\partial v_4} = 1 \quad (22)$$

$$\frac{\partial v_4}{\partial v_3} = \frac{1}{v_3} \quad (23)$$

$$\frac{\partial v_3}{\partial v_1} = \frac{1}{v_2} \quad (24)$$

$$\frac{\partial v_1}{\partial x_1} = \cos(x_1) \quad (25)$$

And finally, evaluating equation (8), as well as substituting our intermediate variables in terms of x_1 , to obtain $\frac{\partial y}{\partial x_1}$:

$$\frac{\partial y}{\partial x_1} = \frac{\cos(x_1)}{v_3 \cdot v_2} \quad (26)$$

$$= \frac{\cos(x_1)}{v_1} \quad (27)$$

$$= \frac{\cos(x_1)}{\sin(x_1)} \quad (28)$$

This process of traversing through the graph backwards is known as reverse accumulation auto differentiation, and in practice each of the gradients would be stored numerically at computation, rather than computing a symbolic version of the final derivative and plugging the values in.

1.2.1.4 Reinforcement Learning (RL)

Reinforcement learning aims to capitalize on learning through an environment (similar to how infants learn by playing), without a teacher, to improve computer's ability at a given task. Simply put, problems in reinforcement learning involve learning how to map certain situations to certain actions to maximize a reward signal. The models are not told what actions to take and instead must explore the environment and available actions to learn what yields the most rewards. In some cases, actions may not only affect immediate rewards, but also the possibility and/or size of future rewards.



Figure 10: Diagram of Markov Decision Process

One challenge that becomes apparent in reinforcement learning methods (supervised and unsupervised), is that of *exploration vs exploitation*. This problem comes from the fact that to maximize any given reward signal, an agent must take the action which gains the most reward i.e. it must *exploit* what it knows, and choose the best action. However, to have knowledge of which action is best at any given state, the agent must have *explored* many different actions - in many different states. This challenge will be addressed in more depth later on, during the design of the agent.

There are two main approaches to reinforcement learning problems, a model-free solution or model solution. Within model-free solutions there is another two categories of policy optimization and Q-learning. In policy optimization, the agent learns a policy which maps states to actions. There are two types of these policies - deterministic and stochastic. A deterministic policy maps without uncertainty i.e. the agent will take the same action given the same state. Stochastic policies on the other hand output a distribution which maps a state to the probability of each action.



Figure 11: Representation of Q-learning process

Q-learning can be represented as a tabular learning method, and works by learning the value of a function - usually $Q(s, a)$ - which represents how successful an action was at a certain state. Each of these values is stored in the 'table' and when the model next observes a given state it consolidates the table and chooses the action which was most successful in its past experience.

For model-based solutions, the models 'know' the rules of the games and learn from planning and construct a functional representation of the environment. This is different to non-model based which aims to learn by trial-and-error and experience. AlphaZero for example would be defined as a model-based agent. To better define this, if the agent is able to forecast the reward of an action given any state, allowing it to plan what actions to take, it is model-based.

Model-based solutions are best when trying to optimize the reward for a task - such as playing chess. Whereas model-free solutions are best for tasks such as self-driving cars (as a model-based approach may run over a pedestrian just to try and complete the journey in the shortest time).

1.2.1.5 AlphaZero

As mentioned before, AlphaZero is model-based - but how exactly it works is unique and does not have a naming convention like Q-learning.

AlphaZero uses a process called Monte Carlo Tree Search (MCTS) to build a tree of future variations of a game. This tree is built depth first - so only begins building a new branch, when the state in the branch currently being explored is terminal. When building this tree, MCTS must select a child node - if all possible new actions have been exhausted. To do this a formula called the upper confidence bound (UCB) is used. The UCB formula can be seen below.

$$UCB = \frac{w_i}{n_i} + C \sqrt{\frac{\ln(N_i)}{n_i}} \quad (29)$$

In this equation, w_i represents the number of wins a node has, n_i represents the number of visits a node has, $N - i$ represents the number of visits the parent node has, and C is a coefficient which controls the

levels of exploration in the MCTS model. Also, after each branch has reached a terminal state the MCTS backpropogates and records how many times each node has been visited, and how many times that node has lead to a winning terminal state. Using this recorded data, the MCTS model can build a distribution which models the probability of winning the game for each action in the current state. This allows the model to 'learn' the rules of the game, as illegal actions will have a probability of 0. This distribution can also be called the MCTS' policy.

To improve this process, AlphaZero uses a neural network to help prune some branches of the tree which can be assumed to lead to a losing state. This allows the MCTS to search only more effective lines of play. The exact algorithm for this can be found in [Identification of Algorithms](#)

1.2.2 Similar Systems

Chinook is a checkers playing computer program that was developed at the University of Alberta between 1989 and 2007. The program utilises an opening book, deep-search algorithms alongside a position evaluation function and a solved endgame database for positions with 8 pieces or fewer. All of Chinook's knowledge has been programmed in by humans, with no use of an artificial intelligence model. Despite this, Chinook managed to beat the checkers world champion at the time of 1995, and after this no longer competed, but instead was tasked with solving checkers. (A solved game is one where the outcome can be predicted from any position) - and this was achieved in 2007.



Figure 13: Chess champion Garry Kasparov loses to Chess engine - 1997

For automatic differentiation, some of the most popular systems include PyTorch's autograd, as well as autodiff by JAX. PyTorch has a very well developed library for deep-learning alongside it's autograd, as well as extensive documentation. Additionally, autograd is more similar to regular python than JAX and may be easier to use when first delving into the area. However, JAX builds on top of autograd and accelerates the linear algebra using an XLA backend. This backend optimizes operations for CPU's GPU's and TPU's as well as using a just-in-time compiler to decrease runtimes.

1.2.3 Interview with third party

For my interview, I reached out to an external person who works in the AI industry. We discussed where my time would be most well spent, some industry standards I should try and implement as well as some expansionary objectives I could try to include to improve my project.



Figure 12: Image of Chinook playing a game of checkers - 1992

Other similar systems include chess engines - of which there is a multitude. Chess engines operate similarly to Chinook, creating trees of possible moves and using a minimax algorithm, as well as alpha-beta pruning to determine the best move. Additionally, there are some chess engines which use reinforcement learning, such as Google DeepMind's AlphaZero engine. This engine has also been trained on games such as Go and Shogi. AlphaZero is trained mainly using self-play, where the AI plays against itself and then learns from those games. In a 100 game match against StockFish (one of the most popular chess engines) AlphaZero won 28 games, lost 0 and drew the remaining 72. This shows the power of reinforcement learning as a tool for board games and its possibility to be used for checkers.

To start with, with the large amount of work that needs to be done, if time begins to run short there are lots of out-of-the-box frameworks to use and implement (such as PyTorch and JAX - mentioned earlier), while still allowing me to finish my project with sufficient difficulty. Secondly, the large proportion of my time should be spent researching and understanding the different solutions available which would best suit my problem, and how they are implemented. This research should include, but not be restricted to: reinforcement learning solutions for other games (such as Go and Chess), different attributes of reinforcement learning models and which environments they are most capable in, as well as research on the overheads of similar models.

It is very important that the correct solution is chosen, so minimal time is spent on implementation. As for possible objectives:

Objective	Reasoning
Leverage GPU cores for linear algebra computations	Allows parallelization and will increase computation speeds
Wrap the console checkers board in a GUI	Much easier for users to use and understand
Use a container such as Docker to produce logs	Much better retracing and debugging capabilities
Evaluate the model against Chinook and human players	Allows for rating of model

1.3 Modelling of the problem

1.3.1 Identification of Objects

1.3.1.1 Checkers

For checkers, we will need two classes, a checkers board/environment and a checkers game. The environment will be used for training the agent and It will require methods that execute an action in the environment, as well as functions to compute the available moves and check for a win. An example of UML for this environment can be seen below.

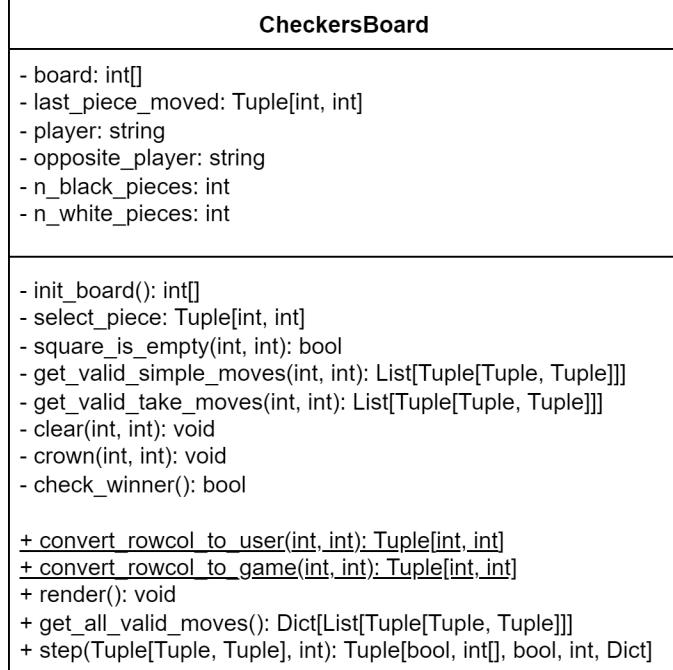


Figure 14: UML for checkers/checkers game

The checkers game varies only slightly and will inherit from the checkers board. It will however add some functionality such as a GUI and the ability to choose different models to play against (such as just the MCTS or the full agent). Additionally, it should allow two people to play each other as checkers usually would.

1.3.1.2 Automatic Differentiation Engine

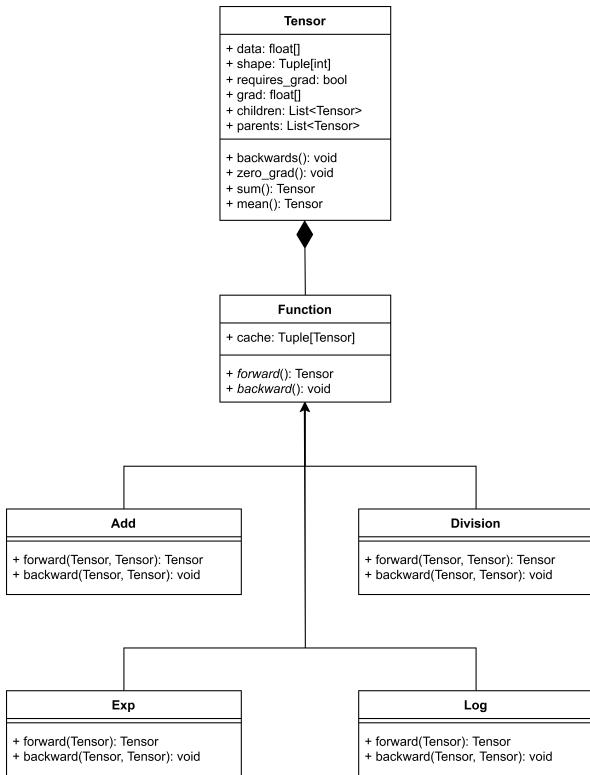


Figure 15: Example UML for Auto Diff

which yields every parameter related to that module. This module class is what all layers will inherit from. It will also have methods to zero out the gradients of all parameters.

Secondly, there will be classes for each different type of layer (dense and activations). The dense layer will have attributes which store the parameters required for computation as well as a method to complete the forward pass through the layer. Activation functions will compute the forward pass, as well as store the gradient function for any tensors which require it. The activation layers will not need any attributes as they have no parameters.

The optimizer class is the base class for any optimizers such as stochastic gradient descent or momentum-based gradient descent. This class will have attributes which store the learning rate as well as any parameters the optimizer is related to. It will also have virtual methods - which will be overridden by the specific optimizer classes - of a step (which updates all parameters), a function to zero the gradient of all parameters and finally a function to add parameters.

Below is an inheritance diagram for the Neural Network Library, without any attributes and methods.

The automatic differentiation engine will be based off of the 'Tensor' object. This object will be similar to a matrix but have some additional attributes, such as its gradient and if its gradient requires computing. If the tensor does require a gradient then it will also store its parent tensors, which will be used for the backward DFS of the computational graph.

Additionally, all operators for the tensor object will have to be overridden to allow for gradient computation if the tensor requires it. The tensor class will also have functions to zero out the gradient as well as begin the backwards traversal through the computational graph. Below is an example of some possible UML for the automatic differentiation library.

Tensor operations will all have a backward and forward method - inherited from the function class. These functions will represent elementary functions (Addition, Division, Log, etc.). The operations such as sum and mean, that operate on the tensor itself will also be represented by a function object.

1.3.1.3 Neural Network Library

There will be a large collection of objects inside of the neural network library, all based around the functionality provided by the tensor object. Firstly, the parameter class will allow for tensors of a given shape to be created and act as parameters for a network - this will also help with backpropagation. The module object will have a parameters property

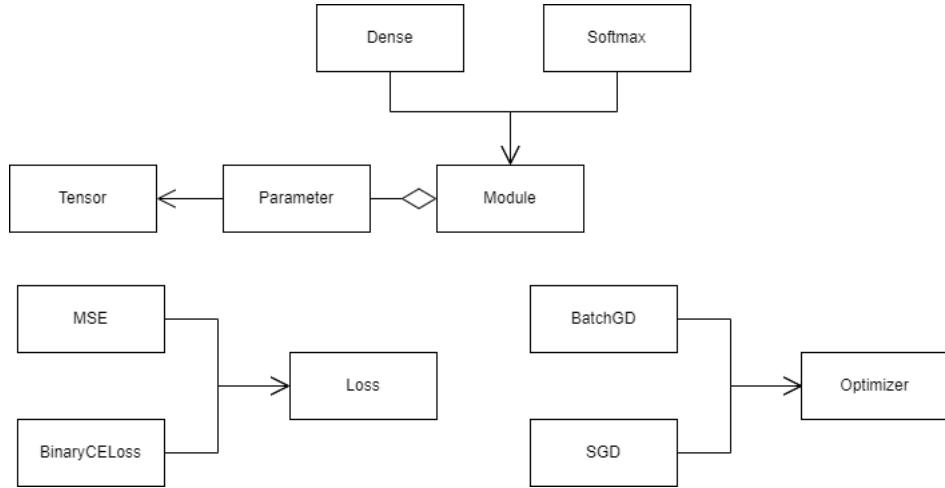


Figure 16: Inheritance diagram for Neural Network Library

1.3.1.4 Reinforcement Learning Solution

The final solution will also require a large set of classes including: a network class - to output the model policy and predicted value; a monte carlo tree search (MCTS) class to compute a search of available actions to train the network towards - composed of a node class to make up this tree - and finally an agent class which abstracts the processes of the previous classes.

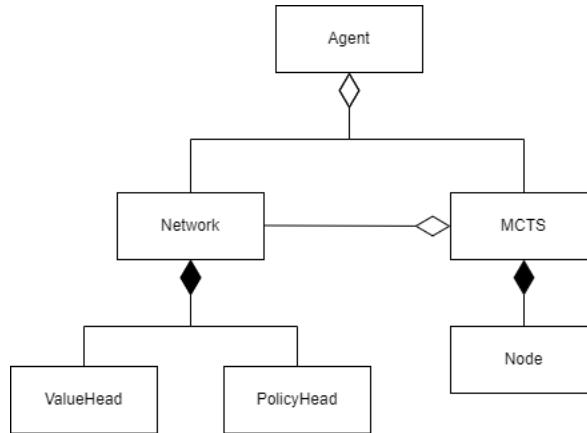


Figure 17: Reinforcement Learning Solution UML

1.3.2 Identification of algorithms

1.3.2.1 Automatic Differentiation Engine

The auto-diff engine will use a recursive DFS of the computational graph to calculate gradients. It uses a DFS as any of the computational graphs that built are directed acyclical graphs, and therefore DFS provides a way to search the whole graph in $O(n)$ time. Some pseudocode for this search can be seen below.

Algorithm 1 Backward Function of Tensor

```
function BACKWARD(Tensor self, Array grad, Tensor out)
Ensure: t.requiresGrad is True                                ▷ t is the tensor being called to start DFS
    if t.Grad is None then
        grad ← [1 1 1 ...] in shape of t.data
        t.Grad += grad
    if t.Operation is not None then
        t.Operation.backward(t.Grad, t)
```

Algorithm 2 General Backward Function of a Tensor Operation

```
function BACKWARD(Operation self, Array grad, Tensor t)
    savedTensors ← self.cache.0
    for tensor in savedTensors do
        if tensor.requiresGrad is True then
            newGrad ←  $\frac{\partial \text{output}}{\partial \text{tensor}}$           ▷ Computed by chain rule, specific for each operation
            if broadcastedDims then
                Remove broadcasted dims from newGrad      ▷ newGrad.shape should equal tensor.shape
                tensor.backward(newGrad, t)
```

1.3.2.2 Neural Network Library

The two main algorithms in the neural network library are the forward pass and the backward pass through a collection of layers. Pseudocode for these can also been seen below.

Algorithm 3 Forward pass through collection of layers

```
function FORWARD(Module model, Tensor input)
    output ← input
    for layer in model.layers do
        output ← layer.forward(output)    ▷ layer's forward function depends on layer type (Dense/Conv)
    return output
```

Algorithm 4 Training process of a module

```
function FIT(Module model, Optimizer optim, Loss l, Tensor x, Tensor y, int epochs = 100, int verbose  
= 1)  
    startTime ← time.now()  
    for epoch in epochs do  
        prediction ← model.predict(x)                                ▷ Get predictions  
        loss ← l(y, prediction)  
        optimizer.zero_grads()                                         ▷ Calculate how far predictions are from true  
        loss.backwards()                                              ▷ Reset gradients of parameters  
        optimizer.step()                                               ▷ Calculate gradients of loss  
        if verbose is 1 then  
            print(epoch, loss)  
        print(time.now() - startTime)                                    ▷ Update parameter values using calculated gradients  
                                                                ▷ Total time elapsed
```

1.3.2.3 Reinforcement Learning Solution

For the reinforcement learning, one of the most important algorithms is the MCTS which is similar to a DFS. This takes all the available moves in the checkers environment, selects one at random and continues this process until a terminal state is reached. After the terminal state is reached, reward values (+1 for win, -1 for loss, 0 for draw) are backpropagated through the nodes. If all possible available moves have been exhausted at a given node then it selects one using an upper confidence bound (UCB) formula. The general algorithm can be seen in pseudocode below (as well as the nodes expand function).

Algorithm 5 MCTS Tree Building Function

```
function BUILDTREE(currentBoardState)  
    self.root ← currentBoardState  
    for search in number_of_searches do  
        node ← self.root                                         ▷ Return to start of tree for each search  
        if node.branchesAvailableToExpand == 0 then  
            node ← node.select_child()  
        while node.terminal is False AND node.branchesAvailableToExpand > 0 do  
            node ← node.expand()                                     ▷ Depth-first expansion until terminal state  
            node.backpropagate(node.reward)
```

The second algorithm is the process of the final agent model selecting its actions and can be seen on the next page.

Algorithm 6 Agent Policy Selection

```
function TRAINAGENT(self, initialNeuralNetwork, hyperParameters)
    nnPrevious ← initialNeuralNetwork
    for mctsEpoch in hyperParameters.mctsEpochs do
        trainingExamplesList ← SPV(hyperParameters.maxExamples)
        for episode in hyperParameters.episodes do
            examples ← list()
            game ← new CheckersGame
            while !game.terminated do
                mcts ← MCTS(nnPrevious, game.currentBoardState, hyperParameters).buildTree()
                if game.step < hyperParameters.Threshold then
                    action ~ mctsPolicy.getPolicy()                                ▷ Select random move
                else
                    action ← mcts.getAction()
                examples.append((game.currentBoardState, mcts.getPolicy(), initialNeuralNetwork.value))
            trainingExamplesList.append(examples)
        for networkEpoch in hyperParameters.networkEpochs do
            batch ← trainingExamplesList.getBatch(hyperParameters.batchSize)
            nnNew ← nnPrevious.train(batch, hyperParameters)
            numberNNNewWins = 0
            for compareGame in hyperParameters.compareGames do
                game ← new CheckersGame
                while !game.terminated do
                    if game.player == player1 then
                        action ← MCTS(nnPrev, game.currentBoardState, hyperParameters).getAction()
                    else if game.player == player2 then
                        action ← MCTS(nnNew, game.currentBoardState, hyperParameters)
                nnNewWinPct = numberNNNewWins/hyperParameters.compareGames
                if nnNewWinPct > hyperParameters.replaceThresholdPct then
                    nnPrev ← nnNew
    return nnPrev
```

1.3.3 Mathematical Formula

1.3.3.1 Automatic Differentiation Engine

In this section I will simply define the formulas for the forward and backward pass of many different functions that will be implemented into the autograd library.

Addition

Forward

$$a + b = y \quad (30)$$

Backward

$$\frac{\partial y}{\partial a} = 1 + 0 \quad (31)$$

$$\frac{\partial y}{\partial b} = 0 + 1 \quad (32)$$

Negation

Forward

$$y = -a \quad (33)$$

Backward

$$\frac{\partial y}{\partial a} = -1 \quad (34)$$

Multiplication

Forward

$$a \cdot b = y \quad (35)$$

Backward

$$\frac{\partial y}{\partial a} = b \quad (36)$$

$$\frac{\partial y}{\partial b} = a \quad (37)$$

Division

Forward

$$\frac{a}{b} = y \quad (38)$$

Backward

$$\frac{\partial y}{\partial a} = \frac{1}{b} \quad (39)$$

$$\frac{\partial y}{\partial b} = -\frac{a}{b^2} \quad (40)$$

Matrix multiplication

Forward

$$AB = Y$$

where A and B are matrices (41)

Backward

$$\frac{\partial Y}{\partial A} = B^T \quad (42)$$

$$\frac{\partial Y}{\partial B} = A^T \quad (43)$$

Powers

Forward

$$a^b = y \quad (44)$$

Backward

$$\frac{\partial y}{\partial a} = b a^{b-1} \quad (45)$$

$$\frac{\partial y}{\partial b} = a^b \ln(a) \quad (46)$$

Mean

Forward

$$\frac{\sum A}{n} = y, \quad A = (x_0, x_1, x_2, \dots, x_n) \quad (47)$$

Backward

$$\frac{\partial y}{\partial A} = B, \quad B = \left(\frac{1}{n_0}, \frac{1}{n_1}, \frac{1}{n_2}, \dots, \frac{1}{n_n} \right) \quad (48)$$

Sum

Forward

$$\sum A = y, \quad A = (x_0, x_1, x_2, \dots, x_n) \quad (49)$$

Backward

$$\frac{\partial y}{\partial A} = B, \quad B = (1_0, 1_1, 1_2, \dots, 1_n) \quad (50)$$

Natural Log

Forward

$$\ln(a) = y \quad (51)$$

Backward

$$\frac{\partial y}{\partial a} = \frac{1}{a} \quad (52)$$

Exponentiation

Forward

$$e^a = y \quad (53)$$

Backward

$$\frac{\partial y}{\partial a} = e^a \quad (54)$$

Conv2D

Forward

$$Y_i = B_i + \sum_{j=1}^n X_j \star K_{ij}, \quad i = 1 \dots d \quad (55)$$

Similar to the current systems, the star denotes the valid cross-correlation of the two matrices.

Example of cross-correlation

Cross correlation can be imagined as placing the kernel ontop of the input and summing the products of the elements which overlap, and then moving the kernel and repeating the process.

$$\begin{bmatrix} 1 & 6 & 2 \\ 5 & 3 & 1 \\ 7 & 0 & 4 \end{bmatrix} \star \begin{bmatrix} 1 & 2 \\ -1 & 0 \end{bmatrix} = \begin{bmatrix} 1 \cdot 1 + 2 \cdot 6 + -1 \cdot 5 + 0 \cdot 3 & 1 \cdot 6 + 2 \cdot 2 + -1 \cdot 3 + 0 \cdot 1 \\ 1 \cdot 5 + 2 \cdot 3 + -1 \cdot 7 + 0 \cdot 0 & 1 \cdot 3 + 2 \cdot 1 + -1 \cdot 0 + 0 \cdot 4 \end{bmatrix} \quad (56)$$

$$= \begin{bmatrix} 8 & 7 \\ 4 & 5 \end{bmatrix} \quad (57)$$

This method of cross-correlation is known as valid cross correlation (where the upper left element of the kernel begins on the upper left element of the input).

Backward

In this backward function, $\frac{\partial Z}{\partial Y}$ will denote the gradient being received from upstream.

$$\frac{\partial Z}{\partial B_i} = \frac{\partial Z}{\partial Y_i} \quad (58)$$

$$\frac{\partial Z}{\partial K_{ij}} = X_j \star \frac{\partial Z}{\partial Y_i} \quad (59)$$

$$\frac{\partial Z}{\partial X_i} = \sum_{j=1}^d \frac{\partial Z}{\partial Y_i} * K_{ij} \quad (60)$$

The asterisk in the final expression represents the full convolution operation, which is the same as a full cross-correlation, but with the kernel rotated 180 degrees. Additionally, a full correlation begins the operation with the bottom left element of the kernel being placed upon the top left element of the input - and then the sliding and summing of products process remains the same.

1.3.3.2 Neural Network Library

Activation Layers

In both dense and convolutional layers, the only operation that is occurring is multiplication and/or matrix multiplication. Graphically, these operations only create linear transformations however in machine learning some problems require non-linearity. To allow for non-linearity neural networks use activation layers which are non-linear functions, normally placed between layers.

Tanh

$$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}} \quad (61)$$

Sigmoid

$$S(x) = \frac{1}{1 + e^{-x}} \quad (62)$$

Softmax

The softmax function is used to turn a vector into a probability distribution. For example, if we had data in the shape (batches, classes) and we wanted to getAction the probability of each batch being in a certain class we would use the softmax function across the axis with the number of class, in this case axis 1.

$$\sigma(x) = \frac{e^x}{\sum e^x} \quad \text{Sum across axis with number of classifications in } x \quad (63)$$

1.3.3.3 Reinforcement Learning Solution

UCB Formula

The UCB formula is used in MCTS to select the most favourable child when the parent node has been fully expanded. The formula can be seen below, where N_i is the number of visits the parent node has, n_i is the number of visits the child node has and w_i is the value count that the node holds. The value count is updated each iteration during backpropagation. Finally, C is a coefficient which determines the levels of exploration and exploitation.

$$UCB(Node_i) = \frac{w_i}{n_i} + C \sqrt{\frac{\ln N_i}{n_i}} \quad (64)$$

1.4 Set of objectives

1.4.1 Checkers Game

1. An easy-to-use and intuitive graphical interface for the checkers game
 - (a) A working version of checkers - obeying all rules stated in [1.2.1.1](#)
 - (b) A small text tutorial on the rules of checkers
 - (c) A visual representation of the checkers board
 - (d) A method of displaying all legal available moves in a selected position
 - (e) Informative error messages if an illegal move is attempted
2. Implementation of various decision-making agents for users to play against
 - (a) Option for the user to choose which agent to play against (MCTS, Full Agent)
 - (b) Ability to adjust hyperparameters of the agents from the checkers game interface (see [RLS Objectives a-e](#))
 - (c) Information printed to the user pertaining to the agent's decision-making progress

1.4.2 Autodifferentiation engine

1. A unit-tested autograd engine capable of calculating derivatives on a variety of functions including:
 - (a) Addition
 - (b) Negation
 - (c) Multiplication
 - (d) Division
 - (e) Matrix Multiplication
 - (f) Powers
 - (g) Mean
 - (h) Sum
 - (i) Natural Log
 - (j) Exponentiation
 - (k) Transpose
 - (l) Padding
 - (m) 2D Convolution

1.4.3 Neural Network Library

1. A modular machine learning library including:
 - (a) The ability to save a model
 - (b) Support for storing a collection of layers/modules
 - (c) Dense/Fully Connected Layer
 - (d) 2D Convolution Layer
 - (e) Reshape Layer
 - (f) Activation Layers:
 - i. ReLU
 - ii. Softmax
 - iii. Sigmoid
 - iv. Tanh
 - (g) Gradient descent optimization
 - (h) Loss Functions:
 - i. Mean Squared error
 - ii. Binary Cross Entropy Loss
 - iii. AlphaZero's Loss Function

1.4.4 Reinforcement Learning Solution

1. A training algorithm with hyperparameters of:
 - (a) Number of example games to be played
 - (b) Number of comparison games to be played
 - (c) Number of searches the MCTS completes
 - (d) The exploration exploitation coefficient for UCB in MCTS
 - (e) The percentage of comparison games the new network must win to replace the previous
 - (f) The number of epochs for neural network training
 - (g) The number of iterations where a new network is trained
2. The functionality to load a saved agent and play against it

2 Design

2.1 File Structure and Project Organisation

```
*ns-*eec-*te-*cg-*rt
└── poetry.lock
└── pyproject.toml
```

At the very top level of my file structure I have a poetry.lock and a pyproject.toml file. These files are automatically generated by a python package called poetry. This package provides exceptionally easy dependency and virtual environment management. Allowing me to setup machines other than my own - to train the model - more easily. The final file inside of this directory actually represents a collection of files - with each * denoting a different value that each of the parameters for the final model can take (with the two letters allowing me to identify said parameter). These files are created using python's built in pickle object serialization library.

```
nea-code
└── AlphaMCTSselfPlay.py
    ├── AlphaModel.py
    ├── CheckersGamePvP.py
    ├── CheckersGuiGame.py
    ├── MCTSselfPlay.py
    ├── PvMCTS.py
    ├── TestNNXOR.py
    └── TrainAgent.py
```

At one level inside of my main code-storing file I have files that run main operations in my code. This includes specific game loops, such as getting the user to play vs MCTS - as well as processes such as beginning agent training and testing of the Neural Network Library using a simple problem such as XOR.

```
nea-code
└── agent
    ├── agent.py
    ├── consts.py
    ├── memory_types.py
    └── __init__.py
```

Inside of nea-code I also create a subdirectory for each module of my project. In the above example this module is the agent class. As well as storing the file which defines the agent class (agent.py) it also contains an __init__.py file (which allows python to recognise the subdirectory of nea-code as a module and defines what dependencies to import when this module is called), a consts.py (which contains values, and functions to create values, which do not change throughout the running of the code) and a memory_types.py file, where dataclasses such as SAP and SPV (which will be defined later on. hereTODO ADD LINK) are defined.

```
nea-code
└── checkers_gui
    ├── buttons.py
    ├── checkers_gui.py
    ├── consts.py
    ├── helpers.py
    └── __init__.py
        └── media
            └── images
                ├── black_k.png
                ├── black_r.png
                ├── checkers_board_no_pieces.png
                ├── white_k.png
                └── white_r.png
```

The next module is the GUI add-on for the checkers game. Similarly to the agent module, it contains

`__init__.py`, `consts.py`, a main class definition file (`checkers_gui.py`) and other class definitions file (`buttons.py`). However, the checkers GUI module also contains a `helpers.py` file - which holds onto functions that are helpful relevant to the module as well as a subtree of `media/images/` which contains the `.png`'s of the different sprites used in the GUI.

```
nea-code
└── console_checkers
    ├── checkers_game.py
    ├── consts.py
    ├── jit_functions.py
    ├── utils.py
    └── __init__.py
```

Inside of the next module (`console_checkers`) we have again an `init`, `consts` and main definition file. On the other hand, instead of a `helpers.py` file there is a `utils.py`. The difference here is that `utils.py` contains utility functions not strictly relevant to the module. For example, this `utils.py` contains a function which clears the console (a process which would not be limited to the `console_checkers` module). Another difference is the presence of the `jit_functions.py` file. Inside of this module there are lots of functions which run only using primitive types, and have a polynomial time complexity. To reduce the impact this has on performance I use a python package called `numba`, which contains a just-in-time compiler (`jit`) to speed up these functions.

```
nea-code
└── mcts
    ├── consts.py
    ├── mcts.py
    └── __init__.py
```

This MCTS module contains definitions of the MCTS classes, again with a `consts` and `init` file.

```
nea-code
└── ml
    ├── __init__.py
    ├── autograd
    │   ├── consts.py
    │   ├── convolve_funcs.py
    │   ├── jit_functions.py
    │   ├── tensor.py
    │   ├── test_autograd.py
    │   ├── helpers.py
    │   └── __init__.py
    └── nn
        ├── layers.py
        ├── optimizers.py
        └── __init__.py
```

In the above snippet I have defined a module named `ml` (for machine learning), which contains two submodules of `autograd` and `nn` - which contain the autodifferentiation engine and neural network library respectively. The `nn` submodule has a similar contents as the modules talked about previously, with `layers.py` defining the main classes, and `optimizers.py` then defining classes of optimizers to be used alongside the defined layers. The `autograd` submodule however has a few differences. It also has a `jit_functions` and `helpers` file but additionally contains a `convolve_funcs.py` and `test_autograd.py` file. These files are unique to the `autograd` module. The `convolve_funcs` has definitions for complex matrix operations such as convolution and cross-correlation, and the `test_autograd` file contains unit tests (run using `pytest`) to evaluate and ensure the accuracy of the autodifferentiation engine.

```
nea-code
└── network
    ├── network.py
    └── __init__.py
```

The final module is `network`, which contains the definitions for compositions of different layers from the `nn` module - which then build to form the model that is contained within the reinforcement learning agent.

2.2 Checkers

The checkers game solution will be split into a console based version (that will be the main way for the agent to interact with the game) and an additional class built ontop of this console solution which provides a GUI.

2.2.1 Data Structures

When defining data structures such as tuples, a comma followed by an ellipsis indicates another item which is of the same type as most recently defined item in the tuple. e.g. `tuple[int,...]` = `tuple[int, int]`

2.2.1.1 Console Solution

Item	Data Structure	Reasoning and Usage
Board	2D Array	Easy and logical to index into specific squares as well as render.
Last Moved Piece	<code>tuple[int,...]</code>	Represents position of the piece most recently moved - easy to reaccess.
Action	<code>tuple[tuple[int,...],...]</code>	First tuple inside of the outer tuple represents the piece being moved, and the second inside tuple represents where it is being moved to which allows for easy indexing into the board array.

For a single game of checkers, an upper bound for the number of unique states that can be presented on the board can be calculated as follows (this is useful for getting a gauge of the size of training sets compared to the possible set of states):

Let $0 \leq w \leq 12$ and $0 \leq b \leq 12$ where $w, b \in \mathbb{Z}$

Then,

$$\binom{32}{w} \quad (65)$$

Is the number of ways to play w white pieces on the board, and

$$\binom{32-w}{b} \quad (66)$$

Is the number of ways to play b black pieces, given w white pieces already occupy a square. And for each piece, it can be a king or a regular piece, so we multiply by a factor of 2^{w+b} . Finally to find the number of legal configurations for all w and b :

$$\sum_{b=0}^{12} \sum_{w=0}^{12} \binom{32}{w} \binom{32-w}{b} 2^{w+b} \approx 2.3 \times 10^{21} \quad (67)$$

Note: Not all of these positions can be reached via a sequence of legal moves - and the actual number (from the paper "Checkers is solved") is $\approx 5 \times 10^{20}$

Additionally, the set of all possible actions can be written as follows:

$$R = \{0, 1, \dots, 7\} \quad (68)$$

$$C = \{A, B, \dots, H\} \quad (69)$$

$$\text{Possible Actions} = \{(x, y), (w, z) \mid x, w \in R \wedge y, z \in C\} \quad (70)$$

2.2.1.2 GUI Add-on

Item	Data Structure	Reasoning and Usage
CheckersGUI	Class	Inherits from console solution and stores information about the game such as the piece the user has selected.
Selected Piece	tuple[int,...]	Represents position of the selected piece - easy to index into board array.
Main Menu	Class	Stores, displays and allows user to change information regarding the parameters and game mode that the user would like to play with. Also initiates game loops.
Button	Class	Provides a simpler to use and understand interface between pygame surfaces and user interactions.
Constants	Class[T]	Stored in classes named relevant to the constant. These classes contain constant attributes of whatever type that constant is defined as (e.g to get the rgb code for red would require COLOURS.RED and would return (255, 0, 0))

2.2.2 Algorithms

The key algorithms for the checkers game are ones which scan the board for all possible moves (divided into two processes, for simple moves and take moves) and the step algorithm which processes an input action for a given board state. Below is pseudocode for the take moves and step functions can be found below.

Algorithm 7 Get Take Moves

```

function GETVALIDTAKEMOVES(CheckersGame self)
    moves ← []
    for row in range(BOARD_SIZE) do
        for col in range(BOARD_SIZE) do
            piece ← self.board[row, col]
            if piece is of self.currentColour then
                for direction in self.piecesLegalDirections(piece) do
                    if self.inBounds(row + 2*direction[row]) and self.inBounds(col + 2*direction[col]) then
                        takingSquare ← (row + direction[row], col + direction[col])
                        moveToSquare ← (row + 2 * direction[row], col + 2 * direction[col])
                        if self.board[moveToSquare] = 0 and self.board[takingSquare] = opponentsPiece then
                            moves.append(((row, col), moveToSquare))
    return moves

```

The above algorithm is for getting valid take moves. The differences between this algorithm and one that gets valid simple moves is that this multiplies all directions by 2 (representing a move over one square and landing on the square beyond that). For getting simple moves, these directions will not be multiplied. Additionally for getting simple moves, the takingSquare variable is not needed - and therefore the check inside the final if statement verifying that the taking square is an opponents piece would also not be needed.

The step algorithm is key to the console checkers solution. It provides a way for the reinforcement learning model to interact by submitting its actions - having them processed with the next state and the reward for the submitted action returned. This process is exactly as defined in the RL Existing Analysis, a markov decision process. In the psuedocode, the return statements have been simplified, but in practice the function would most likely return a tuple containing, at minimum: the next state as a 2D array; whether the game is over or not as a bool; the reward for the submitted action in that state ($r, r \in \mathbb{Z} \cap [-1, 1]$)

Algorithm 8 Checkers Game Step

```

function STEP(CheckersGame self, ACTION action)
    if self.lastPieceMoved is None then
        allMoves ← self.getAllMoves()
        if len(allMoves) = 0 then
            return gameOver, loss
        if action not in allMoves then
            return invalidMove
        else
            self.movePiece()
    else if self.lastPieceMoved is not None then
        allMoves ← self.getTakeMoves()
        if action not in allMoves then
            return invalidMove
        else
            self.movePiece()
    if self.isTakeMove(action) then
        if self.doubleMoveAvailable() then
            lastMovedPiece ← action[pieceMoved]
            self.crownPieces()
        if self.movesWithoutCapture = 40 then
            return draw
        else if self.numberBlackPieces = 1 and self.numberWhitePieces == 1 then
            return draw
        else if self.numberOppositePlayerPieces = 0 then
            return win
        else if self.movesAvailableForOpponent is False then
            return win
        else
            return nextState

```

2.2.3 User Interface

2.2.3.1 Prototyping

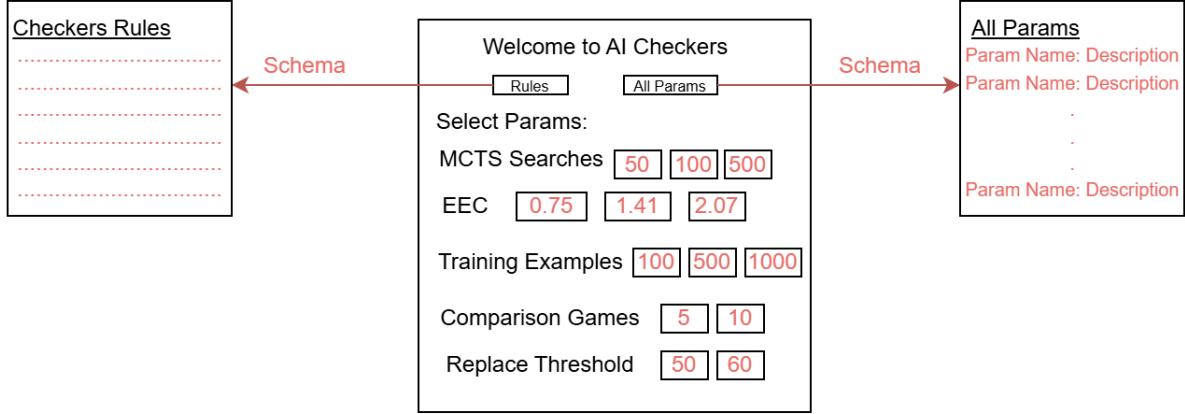


Figure 18: Prototype Main Menu Design

The main menu will provide a starting place for all game loops to begin from - see [2.2.3.2](#) for a visualization. This should allow users to begin a game loop only when the needed parameters for that loop have been selected. For 'User vs User play' therefore this would not require any parameters to be selected. For 'User vs MCTS' play only the 'EEC' and 'MCTS searches' parameters would require values to be selected, but for 'User vs Agent' play it would require a value to be selected for all of the listed parameters on the main menu page. However, this page does not include all the parameters that are needed for each model. Therefore this information could be found inside of the 'All Params' button - with a list of each parameter name, and a description of what it dictates.

2.2.3.2 Structure, Hierarchy and Flow

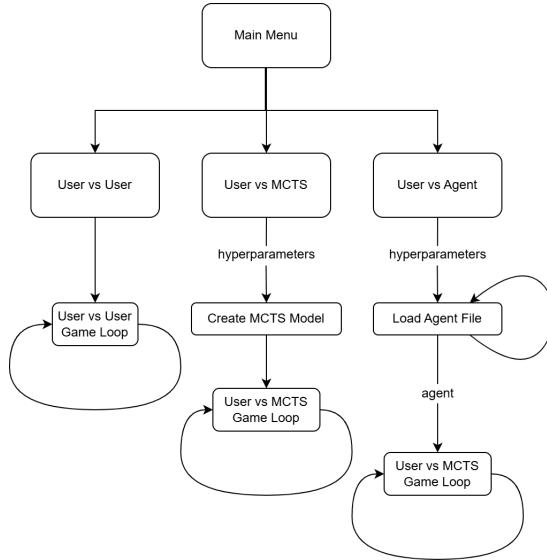


Figure 19: Hierarchy chart for the User Interface

When the user interface is first shown, it displays the main menu. As seen in the most recent section, the main menu will allow users to select various hyperparameters for both the MCTS and complete agent models and then allow the user to begin playing a game of checkers against the desired model. For MCTS,

the model will be instantiated inside of the program, and a tree build will be run on each turn of the MCTS. For the agent objects, its instantiation will be stored inside of a respectively named pickle file - see [2.1](#) - and then loaded into the program.

2.3 Autodifferentiation Engine

The autodifferentiation engine is based mainly around one class - the 'Tensor' class. This class replicates types such as floats, integers, and matrices while adding functionality for automatic differentiation of operations on these types.

2.3.1 Data Structures

Item	Data Structure	Reasoning and Usage
Tensor Data	2D Array	Can represent floats, ints and matrices - with a scalar just being a 1x1 matrix.
Tensor Shape	tuple[int,...]	Represents the shape of the data matrix. Useful for iterating through data.
Tensor Children + Parents	list[Tensor]	Allows us to store the tensors used to create and created by the given tensor - which then builds the function graph. (see example)
Tensor Function Graph	Graph	Represents a graph of operations upon or between input tensors, intermediary tensors, and a final output of a composition of operations. Allows for backward traversal and therefore computation of partial derivatives of the output with respect to its inputs.

2.3.2 Algorithms

The key algorithms for this module include the search backwards through the function graph - which uses a depth first search and various algorithms relating to convolution and correlation operations on matrices, and their derivatives. Some smaller processes include the removing of extra dimensions which are a result of numpy's broadcasting.

2.3.3 Structure, Hierarchy and Flow

2.3.4 Non-standard Design

2.4 Neural Network Library

2.4.1 Data Structures

2.4.2 Algorithms

2.4.3 Structure, Hierarchy and Flow

2.4.4 Non-standard Design

2.5 Reinforcement Learning Solution

2.5.1 Data Structures

2.5.2 Algorithms

2.5.3 Structure, Hierarchy and Flow

2.5.4 Non-standard Design

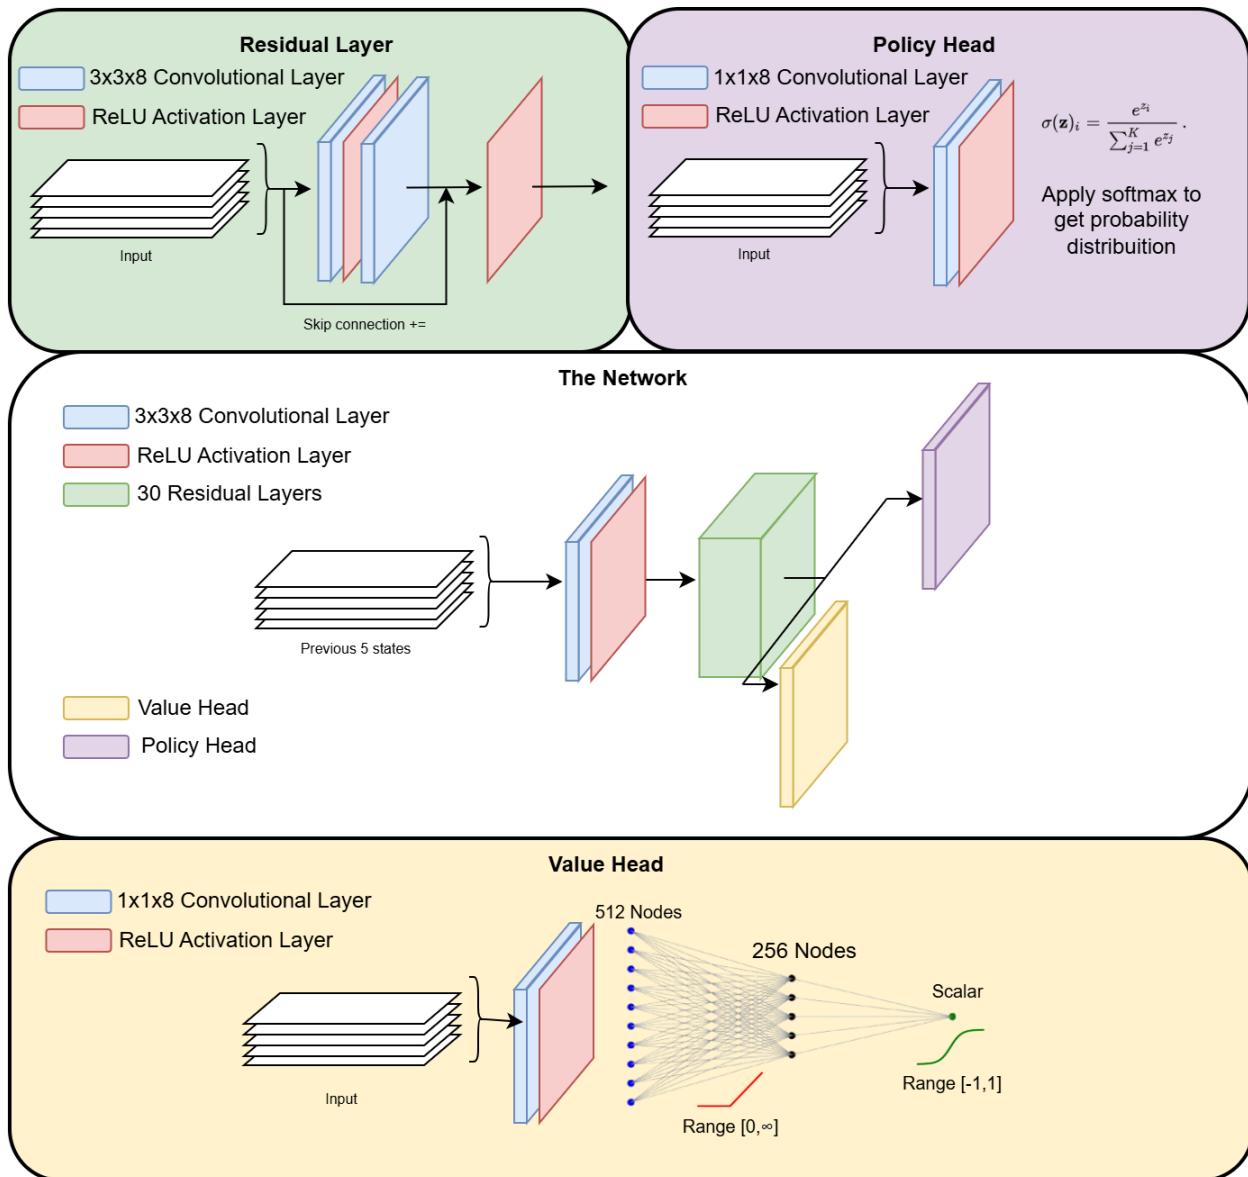


Figure 20: A representation of the AlphaNetwork

3 Research Log

3.1 Checkers

3.1.1 Rules

<https://www.mastersofgames.com/rules/checkers-rules.htm>

3.1.2 Possible States

<https://math.stackexchange.com/questions/2240519/how-to-find-the-amount-of-possible-piece-configuration>

3.2 Neural Networks

<https://www.ibm.com/topics/neural-networks>

3.2.1 Vectorization of summations

https://courses.cs.washington.edu/courses/cse446/20wi/Lecture8/08_Regularization.pdf

3.2.2 Convolutional Networks

<https://www.youtube.com/watch?v=Lakz2MoHy6o>

3.3 Automatic Differentiation

https://en.wikipedia.org/wiki/Automatic_differentiation

3.4 AlphaZero

<https://ai.stackexchange.com/questions/13156/does-alphazero-use-q-learning> <https://suragnair.github.io/posts/alphazero.html> Explanations of AlphaZero's loss function https://www.davidsilver.uk/wp-content/uploads/2020/03/unformatted_final_mastering_go.pdf Methods and Neural Network architecture section

<https://liacs.leidenuniv.nl/~plaata1/papers/CoG2019.pdf> AlphaZero policy explanation

3.5 Existing Auto differentiation systems

<https://medium.com/@utsavstha/jax-vs-pytorch-a-comprehensive-comparison-for-deep-learning-10a84f934e17> PyTorch vs JAX

3.6 Reinforcement Learning

<https://web.stanford.edu/class/psych209/Readings/SuttonBartoIPRLBook2ndEd.pdf> What is reinforcement learning?

<https://smartlabai.medium.com/reinforcement-learning-algorithms-an-intuitive-overview-904e2dff5bbc> Intro to reinforcement learning

<https://medium.com/the-official-integrate-ai-blog/understanding-reinforcement-learning-93d4e34e5698> Different types of reinforcement learning

<https://neptune.ai/blog/model-based-and-model-free-reinforcement-learning-pytennis-case-study> Model-based vs Model-free learning